

Systems Programming

Lecture 8 | 18th of November 2018

Michel Steuwer | <http://michel.steuwer.info> | michel.steuwer@glasgow.ac.uk



Level 3 Student Survey

<http://bit.do/lv3survey>

Recap of last week

- C++ provides **thread** facilities similar to Pthread, but easier to use:

```
auto t = std::thread([](){ printf("Hello World from a thread!\n"); }); t.join();
```

- For this we learned two useful features of C++: type inference with **auto** and **lambda expressions**
- We use the idea of temporary **ownership** of a **std::mutex** to provide mutual exclusion without the possibility to forget to unlock
- To make data structures thread safe we can **encapsulate** an unsafe version and provide a thread safe interface to them
- A **task** is an abstraction of a thread, where we no longer think about low level synchronisation mechanisms, but instead about **communication** between tasks
- **std::async** launches a task and communicates its result back via a **std::future**
- A **std::future** is the reading end of a communication channel. The writing end is called a **std::promise**

Thread safe interfaces

- We discussed last week the encapsulation of a non-thread-safe data structure to make it thread-safe
- This might involve changing the interface of the data structure! Let's have a look at `std::list`:

```
while ( list.size() > 0 ) {  
    auto item = list.front();  
    list.pop_front();  
    // ...  
}
```

- In the code snippet we want to remove elements until the list is empty
- The interface of `std::list` splits this into three functions: `size`, `front` and `pop_front`
- Even if we assume that each individual function is protected with a mutex, there are problems
- For example could two threads simultaneously execute `list.front()`;
both access the first element of the list; and then
both threads call `pop_front()`, resulting in the second element to be removed but never accessed

A thread-safe interface for `std::list` - Take 1

- To address this we can change the interface into a single function:

```
struct maybe_safe_list {  
private: std::list<int> list; std::mutex m;  
public:  
    void push_back(int i) { ... }  
    int size() { ... }  
    int ts_pop_front() {  
        std::unique_lock<std::mutex> l(m); // protect the function body  
        auto item = list.front(); list.pop_front();  
        return item;  
    }  
};
```

- We have to adapt the client code as well:

```
while ( list.size() > 0 ) { auto item = list.ts_pop_front(); /* ... */ }
```

- This resolves the issue, as now `front()` and `pop_front()` have to be executed together
- Is the implementation now thread-safe?

A thread-safe interface for `std::list` - Take 1

- To address this we can change the interface into a single function:

```
struct maybe_safe_list {  
private: std::list<int> list; std::mutex m;  
public:  
    void push_back(int i) { ... }  
    int size() { ... }  
    int ts_pop_front() {  
        std::unique_lock<std::mutex> l(m); // protect the function body  
        auto item = list.front(); list.pop_front();  
        return item;  
    }  
};
```

- We have to adapt the client code as well:

```
while ( list.size() > 0 ) { auto item = list.ts_pop_front(); /* ... */ }
```

- This resolves the issue, as now `front()` and `pop_front()` have to be executed together
- Is the implementation now thread-safe?
- No! Two threads could execute `size()` simultaneously for a list with 1 element;

A thread-safe interface for `std::list` - Take 2

- So we have to combine the size check into the same function:

```
// The function might now return no value when the list is empty
std::optional<int> ts_pop_front() {
    std::unique_lock<std::mutex> l(m); // enter critical section
    if (list.size() > 0) { // check if there is an element to remove
        auto item = list.front();
        list.pop_front();
        return item; // return item after it has been removed
    } else {
        return {}; // nothing to remove; return empty optional
    }
} // exit critical section
```

- This function can now be used in a thread safe way:

```
auto item = list.pop_front();
while ( item.has_value() ) { printf("%d\n", item.value()); item = list.pop_front(); }
```

- We have seen, that it is sometimes required to change the interface design to make an interface thread-safe!


Task management

- Last week we have seen `std::async` as a way to launch *tasks* that communicate their computed results back via an `std::future`
- This way we can write applications without caring about low level synchronisation issues
- In the rest of the lecture we will investigate a task management system
- This is based on the excellent talk "*Better Code: Concurrency*" by Sean Parent
<http://sean-parent.stlab.cc/papers-and-presentations/#better-code-concurrency>
- Sean Parent has more excellent tasks about modern C++ on his website

A basic task system

- We start by building a task system that operates with a single queue to manage the incoming tasks
- Multiple threads will take tasks out of the queue and execute them concurrently
- This follows the guidelines of an multi-threaded programming guide published in 2010 by Oracle

(<https://docs.oracle.com/cd/E19253-01/816-5137/ggedn/index.html>)



std::function for passing functions as arguments

- Each task is defined by the function it executes
- To add a task into our task system we provide a function as argument
- In C we have used function pointers for this:

```
void foo() { ... }  
  
void async( void(*f)() ) { ...}  
int main() { async(foo); }
```

- In C++ we use std::function:

```
void foo() { ... }  
  
void async( std::function<void()> f ) { ... }  
int main() { async(foo); }
```

- We specify the type signature of the function in the < angle brackets >: std::function<void()>

The basic queue implementation

```
struct notification_queue {
private:
    std::list<std::function<void()>> q;    // this is our list of tasks; each task is a function
    std::mutex m;                        // a mutex to protect the queue
    std::condition_variable ready;       // a condition variable indicating available tasks
public:
    std::function<void()> pop() {         // pop removes a task from the queue
        std::unique_lock<std::mutex> lock(m); // enter a critical section to protect the std::list
        ready.wait(lock, [this]{ return !q.empty(); }); // it waits until a task is in the queue
        auto f = q.front(); // access the first task ...
        q.pop_front();      // ... remove it from the queue
        return f;           // ... and return it
    } // exit critical section

    void push(std::function<void()> f) { // push adds a task to the queue
        {
            std::unique_lock<std::mutex> lock(m); // enter critical section to protect the std::list
            q.push_back(f);                      // add task f to the queue
        } // exit critical section
        ready.notify_one(); // notify a waiting thread that a task is in the queue
    }
};
```

The basic task system implementation

```
struct task_system {
private:
    const int count = std::thread::hardware_concurrency(); // init with number of cores
    std::vector<std::thread> threads; // a vector to store all the worker threads
    notification_queue q; // the queue to store the tasks

    void run(int i) { // function executed by each worker thread
        while (true) { auto f = q.pop(); f(); } // forever, we remove a task and execute it
    }
public:
    task_system() { // This is the constructor
        printf("Start task system with %d threads\n", count);
        for (auto n = 0; n != count; n++) { // for each hardware core ...
            threads.emplace_back([this, n]() { run(n); }); // ... we create a thread executing run
        }

        ~task_system() { // This is the destructor
            for (auto n = 0; n != count; n++) { threads[n].join(); } // we ensure to wait for all threads
        }

        void async(std::function<void()> f) { // This is the interface to add tasks to the system
            q.push(f); // we simply push the task into the queue, which will notify a waiting thread
        }
    };
};
```

Stopping Tasks

```
struct notification_queue {
private:
    std::list<std::function<void()>> q; std::mutex m; std::condition_variable ready;
    bool done = false; // this flag indicates the end of executing tasks
public:
    // We change the return type, as pop might now return 'nothing' when done is set
    std::optional<std::function<void()>> pop() {
        std::unique_lock<std::mutex> lock(m);
        ready.wait(lock, [this]{ return (!q.empty()) || done; }); // if done is set we will not wait
        if (q.empty()) return {}; // if the queue is empty (because done was set) we return 'nothing'
        auto f = q.front();
        q.pop_front();
        return f;
    }
    void push(std::function<void()> f) { ... }
    void setDone() {
        {
            std::unique_lock<std::mutex> lock(m); // enter critical region to protect 'done'
            done = true; // set done
        }
        ready.notify_all(); // Notify all threads that they should finish
    }
};
```

Stopping Tasks


```
struct task_system {
private:
    const int count = ... ; std::vector<std::thread> threads; notification_queue q;
    void run(int i) {
        while (true) {
            auto optional_f = q.pop();           // we now return an optional
            if (!optional_f.has_value()) return; // if 'nothing' was returned we exit the function
            auto f = optional_f.value();         // if something was returned we get the task out ...
            f();                                 // ... and execute it
        }
    }
}

public:
    task_system() { ... }

    ~task_system() {
        q.setDone(); // when we destroy the task system, let all threads now they should finish ...
        for (auto n = 0; n != count; n++) { threads[n].join(); } // ... and then wait for them
    }

    void async(std::function<void()> f) { ... }
};
```

Performance of the simple task system



Pretty bad performance, as all threads access the same queue there is a lot of blocking and waiting

Take 2: Improved task system - Multiple Queues

- **Idea:** to reduce blocking and waiting time use multiple queues
- A scheduling policy will decide in which queue to enter the task
- Each thread will exclusively access tasks from their own queue
- Different scheduling policies could be implemented, we pick a very simple, but widely used one: *round robin*
 - Every time a task has been added to a queue we add the next task to the next queue
 - Once we arrive at the last queue, we start again with the first queue
- For this implementation we use the exact same queue implementation as before , we only change the task system

The multi-queue task system implementation

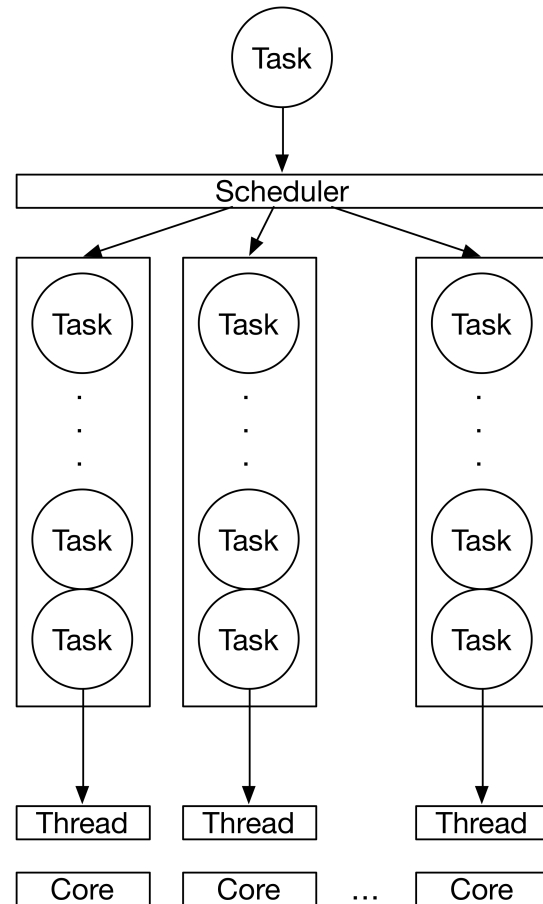
```
struct task_system {
private: const int count = ...; std::vector<std::thread> threads;
    // we now have one queue per thread
    std::vector<notification_queue> q = std::vector<notification_queue>(count);
    // this 'atomic' int can be incremented thread-safely without requiring a lock
    std::atomic<int> index = 0;

    void run(int i) {
        while (true) {
            auto optional_f = q[i].pop(); // now every thread pops the tasks from their queue
            if (!optional_f.has_value()) return;
            auto f = optional_f.value(); f(); } }
public:
    task_system() { ... }

    ~task_system() {
        for (auto n = 0; n != count; n++) { q[n].setDone(); } // indicate each queue to finish
        for (auto n = 0; n != count; n++) { threads[n].join(); }
    }

    void async(std::function<void()> f) {
        // This implements the round robin scheduling
        auto i = index++; // every time increment the index by one
        q[i % count].push(f); // the modulo makes sure that we are staying in bound
    }
};
```

Performance of multi-queue task system



Better performance, but still not great. Some threads might have more work than others.

Take 3: Improved task system - Work stealing

- **Idea:** Let a thread try to help others by stealing a task - if one is available - from the other threads queues. If no task can be stolen, wait on your own queue for a task to arrive.
- Work stealing only works efficiently if we can avoid blocking (i.e. waiting for a task to arrive) when we steal
- In C++ we can *try to acquire* a mutex.
If the mutex is locked we do not block:

```
void try_lock() {  
    std::unique_lock<std::mutex> lock(m, std::try_to_lock);  
    if (!lock) { // if not locked successfully  
        return; } // do not enter critical section  
    // otherwise enter critical section  
    // ...  
}
```

Updated Queue - try_push, try_pop

```
struct notification_queue {
private:
    std::list<std::function<void()>> q; bool done=false; std::mutex m; std::condition_variable ready;
public:
    // new function: tries to pop, if queue is locked it will return 'nothing'
    std::optional<std::function<void()>> try_pop() {
        std::unique_lock<std::mutex> lock(m, std::try_to_lock); // do not wait if the mutex is locked
        if (!lock || q.empty()) return {}; // we have to(!) check if the lock was acquired
        auto f = q.front(); q.pop_front(); return f;
    }
    // new function: tries to push, if queue is locked it will return false
    bool try_push(std::function<void()> f) {
        { std::unique_lock<std::mutex> lock(m, std::try_to_lock);
            if (!lock) return false; // return false to indicate that the push has not been successful
            q.push_back(f);
        }
        ready.notify_one();
        return true; // return true if the task was successfully added to the queue
    }

    std::optional<std::function<void()>> pop() { ... }
    void push(std::function<void()> f) { ... }
    void setDone() { ... }
```

Updated task system - run

```
struct task_system { private: // ...

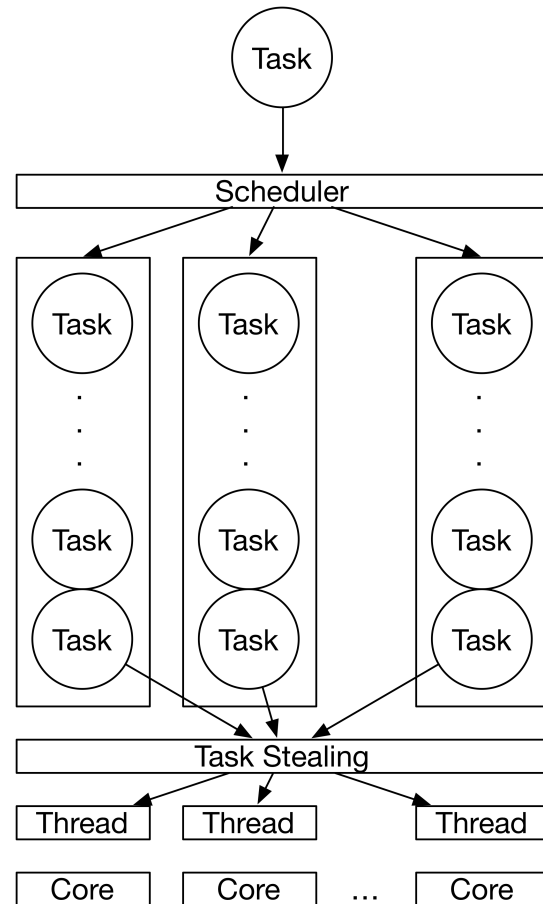
void run(int i) { while (true) {
    std::optional<std::function<void()>> optional_f;
    // first try to steal from all other queues
    for (auto n = 0; n != count; n++) {
        optional_f = q[ (i + n) % count ].try_pop();
        if (optional_f.has_value()) break; // once a steal was successful -> stop looking
    }
    // if we have not stole successfully, we are accessing our own queue ...
    if (!optional_f.has_value()) {
        // ... and potentially block and wait
        optional_f = q[i].pop();
        if (!optional_f.has_value()) return; // if our queue returns no task, we are done
    }
    // if we (somehow) got a task ...
    auto f = optional_f.value(); // ... access ...
    f();                        // ... and execute it
} }

public:
    task_system() { ... }      ~task_system() { ... }
    void async(std::function<void()> f) { ... } // see next slide
};
```

Updated task system - async

```
struct task_system {  
private: // ...  
    void run(int i) { ... }  
  
public:  
    task_system() { ... }  
    ~task_system() { ... }  
  
    void async(std::function<void()> f) {  
        auto i = index++;  
        // first try to insert in all other queues  
        for (auto n = 0; n != count; ++n) {  
            // if an insertion is successful, return from the function  
            if (q[ (i + n) % count ].try_push(f)) return;  
        }  
        // only if no other insertion was successful, add in our own queue  
        // and potentially block  
        q[i % count].push(f);  
    }  
};
```


Performance of work stealing task system



Very good performance.
Work stealing is a crucial technique for a
high-performance task system

System Programming Recap - Part 1

- In this course we have studied systems programming. We have learned, that:
 - **data types** give bit representations in memory a meaning
 - **structs** allow us to implement custom data structures (like link lists or trees)
 - every variable is stored at a fix location **memory**
 - a **pointer** is a variable storing the address of a memory location
 - memory on the **stack** is automatically managed but memory on the **heap** must be managed manually
 - to organise resource and memory management we should think about **ownership**
 - in C++ ownership is implemented by tying the resource management to the **lifetime** of a stack variable
 - there exist a set of **smart pointers** and **containers** for easy and non-leaking memory management
 - **debuggers, static** and **dynamic analysis tools** help to detect and fix bugs

System Programming Recap - Part 2


- In the second part of the course, we talked about programming concurrent systems. We learned, that:
 - **concurrency** is a programming paradigm to *deal with lots of things at once*
 - **parallelism** is about *doing lots of things at once* and this way making a program run faster
 - multiple **threads** share the same address space of a single **process**
 - Pthreads is a threading implementation in C, C++ provides threads in its standard library
 - **mutual exclusion** ensures that two threads don't simultaneously enter their **critical section**
 - **condition variables** are a synchronisation mechanism to avoid busy waiting for a condition
 - the bounded buffer is an example of a **monitor**, an object that encapsulates low level synchronisation
 - **tasks** provide an abstraction over threads:
we think about communicating tasks rather than threads modifying shared state
 - a **task system** allows to efficiently use hardware threads using multiple queues and **work-stealing**

Where to go from here?


- There are multiple *systems* courses that build up on the ideas you have learned in this course:
 - *Operating Systems*
 - *Networked Systems*
 - *Distributed and Parallel Technologies*
 - *Advanced Systems Programming*
- For all of these it is important to have a good understanding of **memory** as discussed in this course
- In *Advanced Systems Programming* you are going to learn *Rust* - a modern systems programming language - that enforces the idea of ownership to guarantee concurrent data race free programs

More Reading and Viewing Tips

- You have seen pretty much everything there is too see of C in this course
- You have pretty much seen only a very small fraction of C++ in this course
- There is much more to discover in C++. These two books are good intermediate reading material:



Great best practice guide of C++



Great best practice guide of modern C++.
Complementary to the first book

- There are many great talks about C++ on Youtube (and other video platforms):
 - *Sean Parents* series on *Better Code*: <http://sean-parent.stlab.cc/papers-and-presentations/>

Next week

- Guest lecture:
Alastair Murray from **Codeplay** will be talking about his industry experience in systems programming
- Codeplay is a Edinburgh based company working on *compilers, debuggers, runtimes, testing systems*, and other software *development tools*.
- They have had (and currently have) Glasgow University students for placements and internships



Systems Programming

Lecture 8 | 18th of November 2018

Michel Steuwer | <http://michel.steuwer.info> | michel.steuwer@glasgow.ac.uk

